

Hierarchical Clustering On Mall Customers data

Importing the libraries

In [2]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import scipy.cluster.hierarchy as sch
6 from sklearn.cluster import AgglomerativeClustering
7
8 import warnings
9 warnings.filterwarnings("ignore")
10 %matplotlib inline
```

Importing the dataset

```
In [3]: 1 dataset = pd.read_csv('C:/Users/Eric/Documents/Jupyter Notebook/practical/data/Mall_Customers.csv')
        2 dataset.head(10)
```

Out[3]:

	CustomerID	Genre	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40
5	6	Female	22	17	76
6	7	Female	35	18	6
7	8	Female	23	18	94
8	9	Male	64	19	3
9	10	Female	30	19	72

```
In [5]: 1 dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                            200 non-null   int64
1   Genre                                 200 non-null   object
2   Age                                   200 non-null   int64
3   Annual Income (k$)                    200 non-null   int64
4   Spending Score (1-100)                 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
In [6]: 1 dataset.isnull().sum()
```

```
Out[6]: CustomerID          0  
Genre          0  
Age            0  
Annual Income (k$)  0  
Spending Score (1-100)  0  
dtype: int64
```

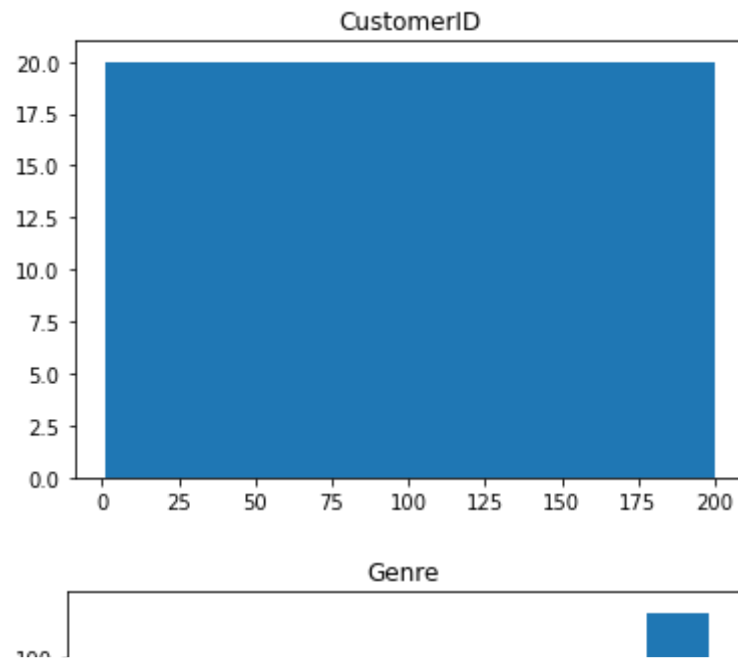
```
In [7]: 1 dataset.describe()
```

```
Out[7]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

EDA

```
In [8]: 1 # df_num = train[['Age', 'SibSp', 'Parch', 'Fare']]
2 for i in dataset.columns:
3     plt.hist(dataset[i])
4     plt.title(i)
5     plt.show()
```




```
In [9]: 1 dataset[['Genre', 'Age']].groupby(['Genre'], as_index=False).mean().sort_values(by='Age', ascending=False)
```

Out[9]:

	Genre	Age
1	Male	39.806818
0	Female	38.098214

```
In [10]: 1 plt.figure(figsize=(8, 5))
        2 sns.barplot(x='Genre', y='Age', data=dataset)
```

Out[10]: <AxesSubplot:xlabel='Genre', ylabel='Age'>




```
In [11]: 1 dataset[['Genre', 'Annual Income (k$)']].groupby(['Genre'], as_index=False).mean().sort_values(by='Annual Income
```

Out[11]:

	Genre	Annual Income (k\$)
1	Male	62.227273
0	Female	59.250000

```
In [12]: 1 plt.figure(figsize=(8, 5))
          2 sns.barplot(x='Genre', y='Annual Income (k$)', data=dataset)
```

Out[12]: <AxesSubplot:xlabel='Genre', ylabel='Annual Income (k\$)'>




```
In [13]: 1 dataset[['Genre', 'Spending Score (1-100)']].groupby(['Genre'], as_index=False).mean().sort_values(by='Spending S
```

Out[13]:

	Genre	Spending Score (1-100)
0	Female	51.526786
1	Male	48.511364

```
In [14]: 1 plt.figure(figsize=(8, 5))
          2 sns.barplot(x='Genre', y='Spending Score (1-100)', data=dataset)
```

```
Out[14]: <AxesSubplot:xlabel='Genre', ylabel='Spending Score (1-100)'>
```




```
In [15]: 1 dataset[['Age', 'Annual Income (k$)']].groupby(['Age'], as_index=False).mean().sort_values(by='Annual Income (k$)')
```

Out[15]:


	Age	Annual Income (k\$)
23	41	101.000000
14	32	87.181818
10	28	85.250000
18	36	81.000000
15	33	80.333333
38	56	79.000000
16	34	79.000000
12	30	76.142857
26	44	75.500000
20	38	74.500000
21	39	72.666667


```
In [16]: 1 plt.figure(figsize=(15, 7))
        2 sns.barplot(x='Age', y='Annual Income (k$)', data=dataset)
```

Out[16]: <AxesSubplot:xlabel='Age', ylabel='Annual Income (k\$)'\>



```
In [17]: 1 dataset.plot(kind="scatter", x="Age", y="Annual Income (k$)")  
2 plt.show()
```




In [18]: 1 dataset[['Age', 'Spending Score (1-100)']].groupby(['Age'], as_index=False).mean().sort_values(by='Spending Score

Out[18]:


	Age	Spending Score (1-100)
21	39	84.666667
12	30	80.285714
11	29	76.600000
6	24	71.500000
4	22	70.000000
10	28	70.000000
3	21	66.800000
14	32	66.000000
17	35	63.888889
13	31	63.875000
5	23	63.333333

```
In [19]: 1 plt.figure(figsize=(15, 7))
        2 sns.barplot(x='Age', y='Spending Score (1-100)', data=dataset)
```

```
Out[19]: <AxesSubplot:xlabel='Age', ylabel='Spending Score (1-100)'>
```




```
In [20]: 1 dataset.plot(kind="scatter", x="Age", y="Spending Score (1-100)")  
        2 plt.show()
```




```
In [21]: 1 plt.figure(figsize=(15, 7))
          2 sns.barplot(x='Annual Income (k$)', y='Spending Score (1-100)', data=dataset)
```

Out[21]: <AxesSubplot:xlabel='Annual Income (k\$)', ylabel='Spending Score (1-100)'\>



```
In [22]: 1 dataset.plot(kind="scatter", x="Annual Income (k$)", y="Spending Score (1-100)")  
        2 plt.show()
```



In [4]:


```
1 # X = dataset.iloc[:, [3, 4]]
2 # X.head()
3 X = dataset.iloc[:, [3, 4]].values
4 X
```

```
Out[4]: array([[ 15, 39],  
               [ 15, 81],  
               [ 16, 6],  
               [ 16, 77],  
               [ 17, 40],  
               [ 17, 76],  
               [ 18, 6],  
               [ 18, 94],  
               [ 19, 3],  
               [ 19, 72],  
               [ 19, 14],  
               [ 19, 99],  
               [ 20, 15],  
               [ 20, 77],  
               [ 20, 13],  
               [ 20, 79],  
               [ 21, 35],  
               [ 21, 66],  
               [ 23, 29],  
               ...])
```


Using the dendrogram to find the optimal number of clusters


In [5]:

```
1 # create dendrogram
2 plt.figure(figsize=(20, 7))
3 dendrogram = sch.dendrogram(sch.linkage(X, method = 'ward'))
4 plt.title('Dendrogram')
5 plt.xlabel('Customers')
6 plt.ylabel('ward distances')
7 plt.show()
```



```
In [25]: 1 plt.figure(figsize=(10, 7))
2         plt.title("Dendrograms")
3         dendrogram = sch.dendrogram(sch.linkage(X, method = 'ward'))
4         plt.axhline(y=200, color='r', linestyle='--')
```

Out[25]: <matplotlib.lines.Line2D at 0x24e024f3640>



Training the Hierarchical Clustering model on the dataset

```
In [26]: 1 # create clusters
          2 hc = AgglomerativeClustering(n_clusters = 5, affinity = 'euclidean', linkage = 'ward')
          3 y_hc = hc.fit_predict(X)
```

```
In [ ]:
```

```
In [27]: 1 print(y_hc)
```

```
[4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4 3 4
 3 4 3 4 3 4 1 4 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 1 1 1 1 1 1 1 1 1 1 1 2 1 2 1 2 0 2 0 2 1 2 0 2 0 2 0 2 1 2
 0 2 0 2 0 2 0 2 0 2 0 2 1 2 0 2 0 2 0 2 0 2 0 2 0 2 0 2 0 2
 2 0 2 0 2 0 2 0 2 0 2 0 2]
```


Visualising the clusters

In [28]:

```
1 plt.scatter(X[y_hc == 0, 0], X[y_hc == 0, 1], s = 100, c = 'red', label = 'Cluster 1')
2 plt.scatter(X[y_hc == 1, 0], X[y_hc == 1, 1], s = 100, c = 'blue', label = 'Cluster 2')
3 plt.scatter(X[y_hc == 2, 0], X[y_hc == 2, 1], s = 100, c = 'green', label = 'Cluster 3')
4 plt.scatter(X[y_hc == 3, 0], X[y_hc == 3, 1], s = 100, c = 'cyan', label = 'Cluster 4')
5 plt.scatter(X[y_hc == 4, 0], X[y_hc == 4, 1], s = 100, c = 'magenta', label = 'Cluster 5')
6 plt.title('Clusters of customers')
7 plt.xlabel('Annual Income (k$)')
8 plt.ylabel('Spending Score (1-100)')
9 plt.legend()
10 plt.show()
```



```
In [29]: 1 plt.figure(figsize=(10, 7))
2 plt.scatter(X[:,0], X[:,1], c=hc.labels_, cmap='rainbow')
3 plt.title('Clusters of customers')
4 plt.xlabel('Annual Income (k$)')
5 plt.ylabel('Spending Score (1-100)')
6 plt.show()
```



```
In [ ]: 1
```

